



# CE/CZ2005: Operating Systems – Lab Experiment 4



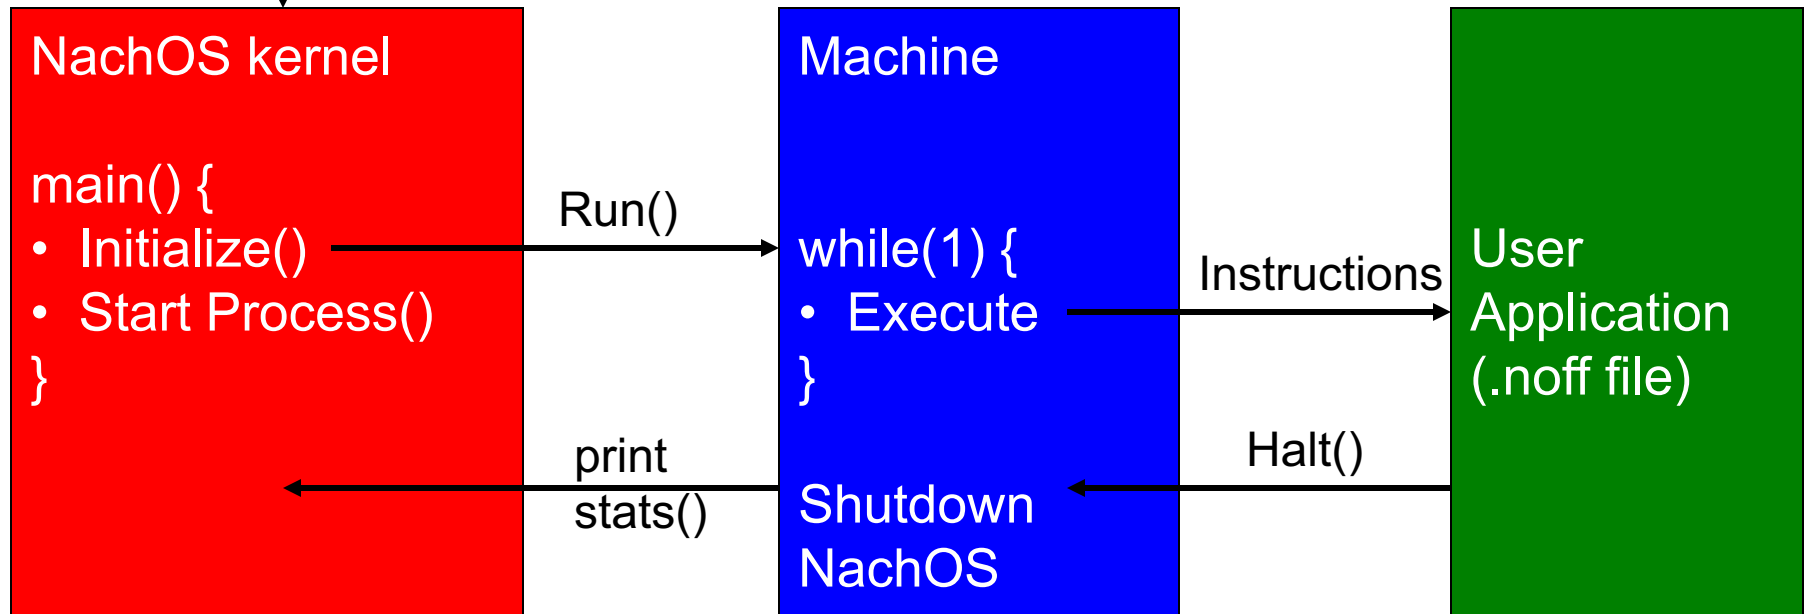
# Outline

- Address translation in NachOS
- Discussion of Experiment 4

# NachOS – Remember how it works?

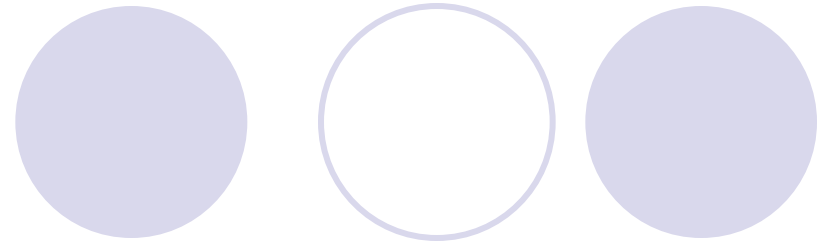
Start NachOS binary  
(./nachos)

-x ../test/vmtest.noff



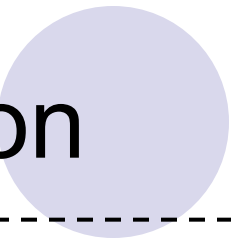
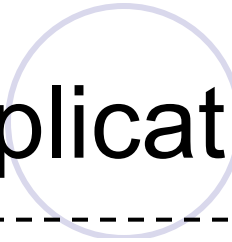
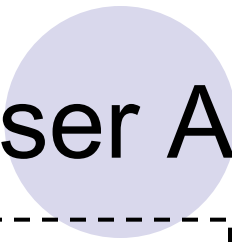
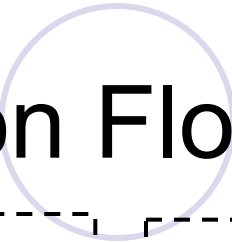
Return to host  
OS shell

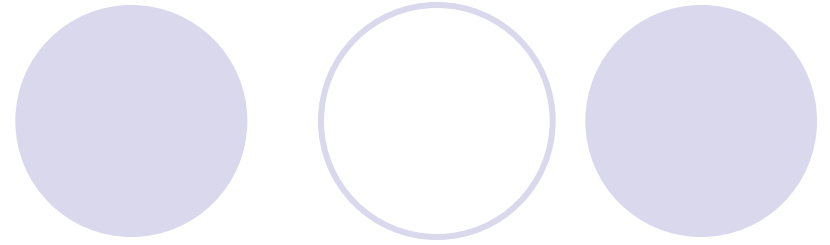
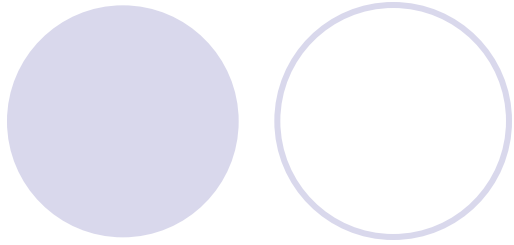
# User Application



- Source code for the user application (`test/vmtest.noff`) can be found in `test/vmtest.c`.
- `vmtest` **executes** 1 child processes (`test/vm.noff`).

# Executive



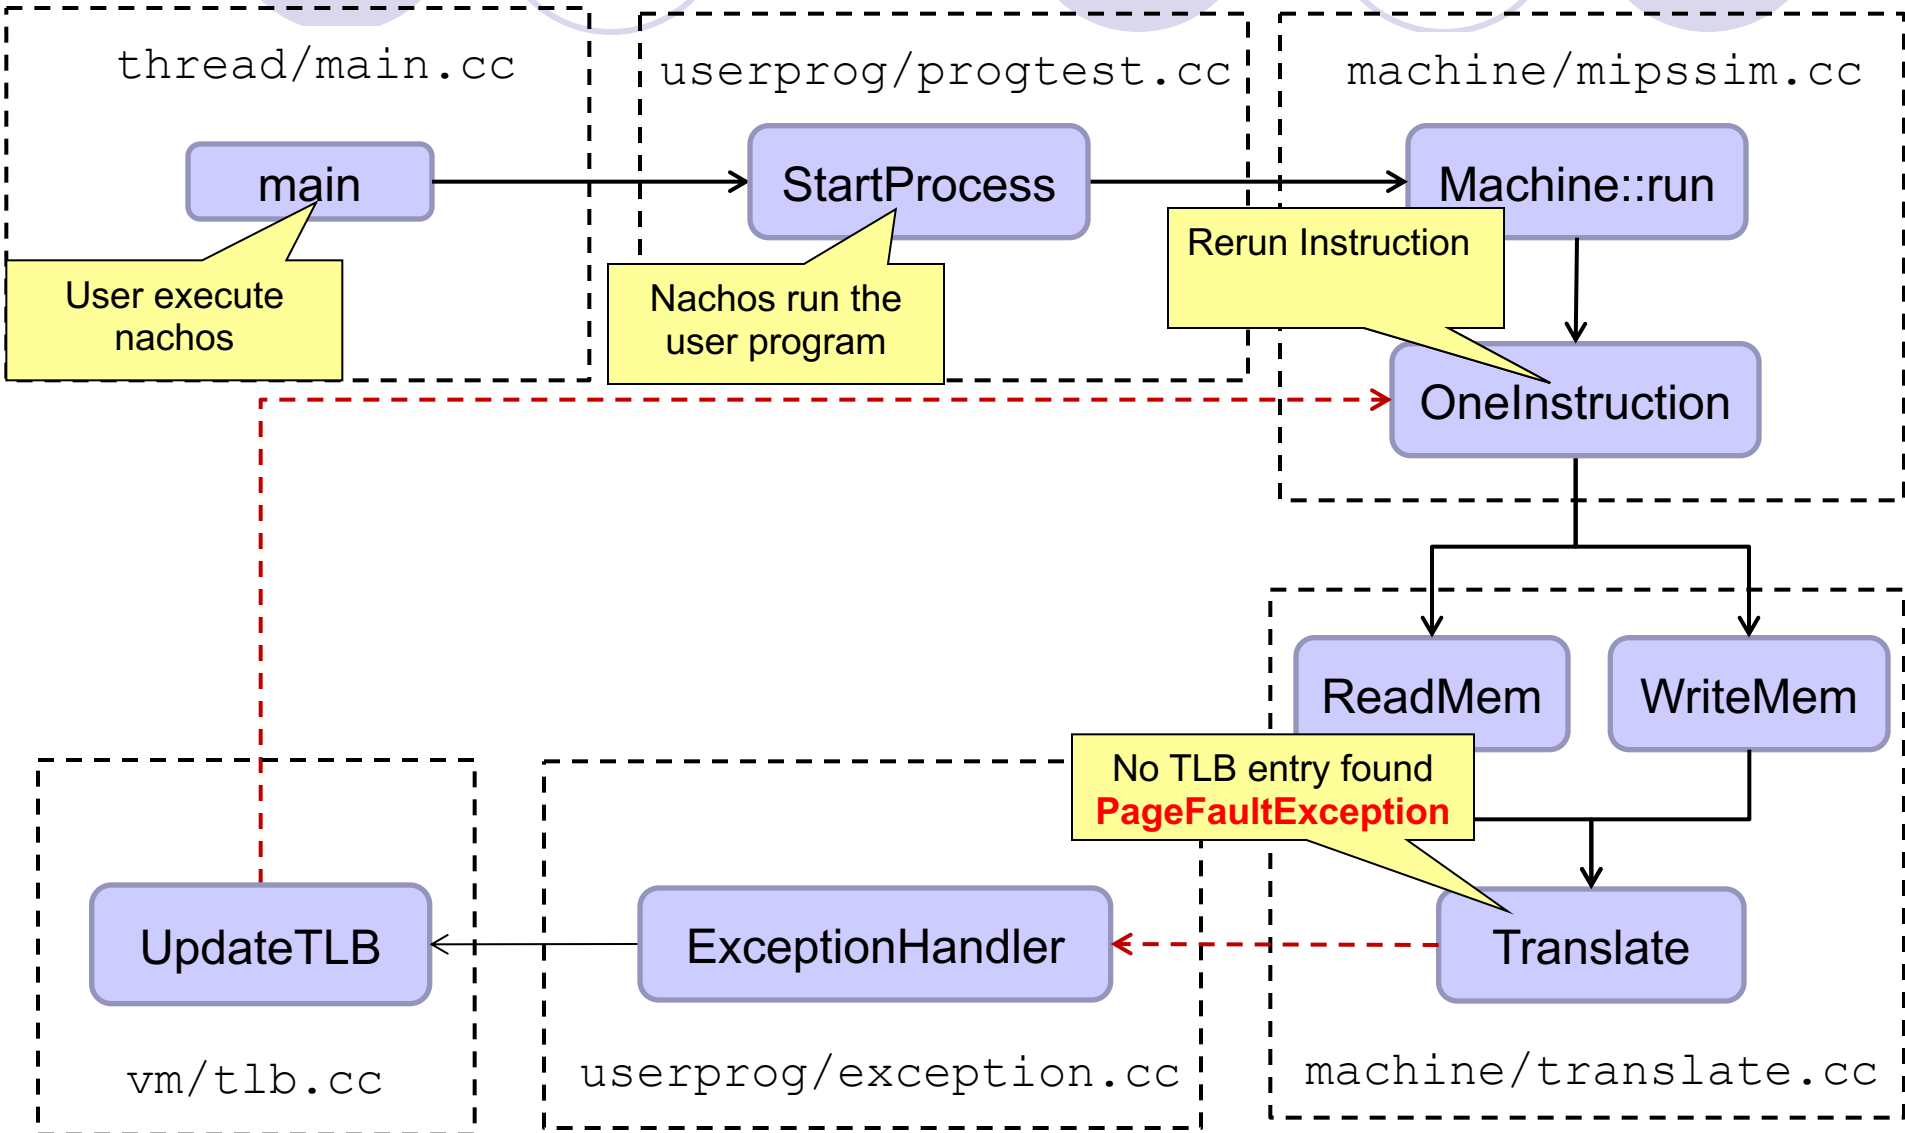


# Address Translation in NachOS

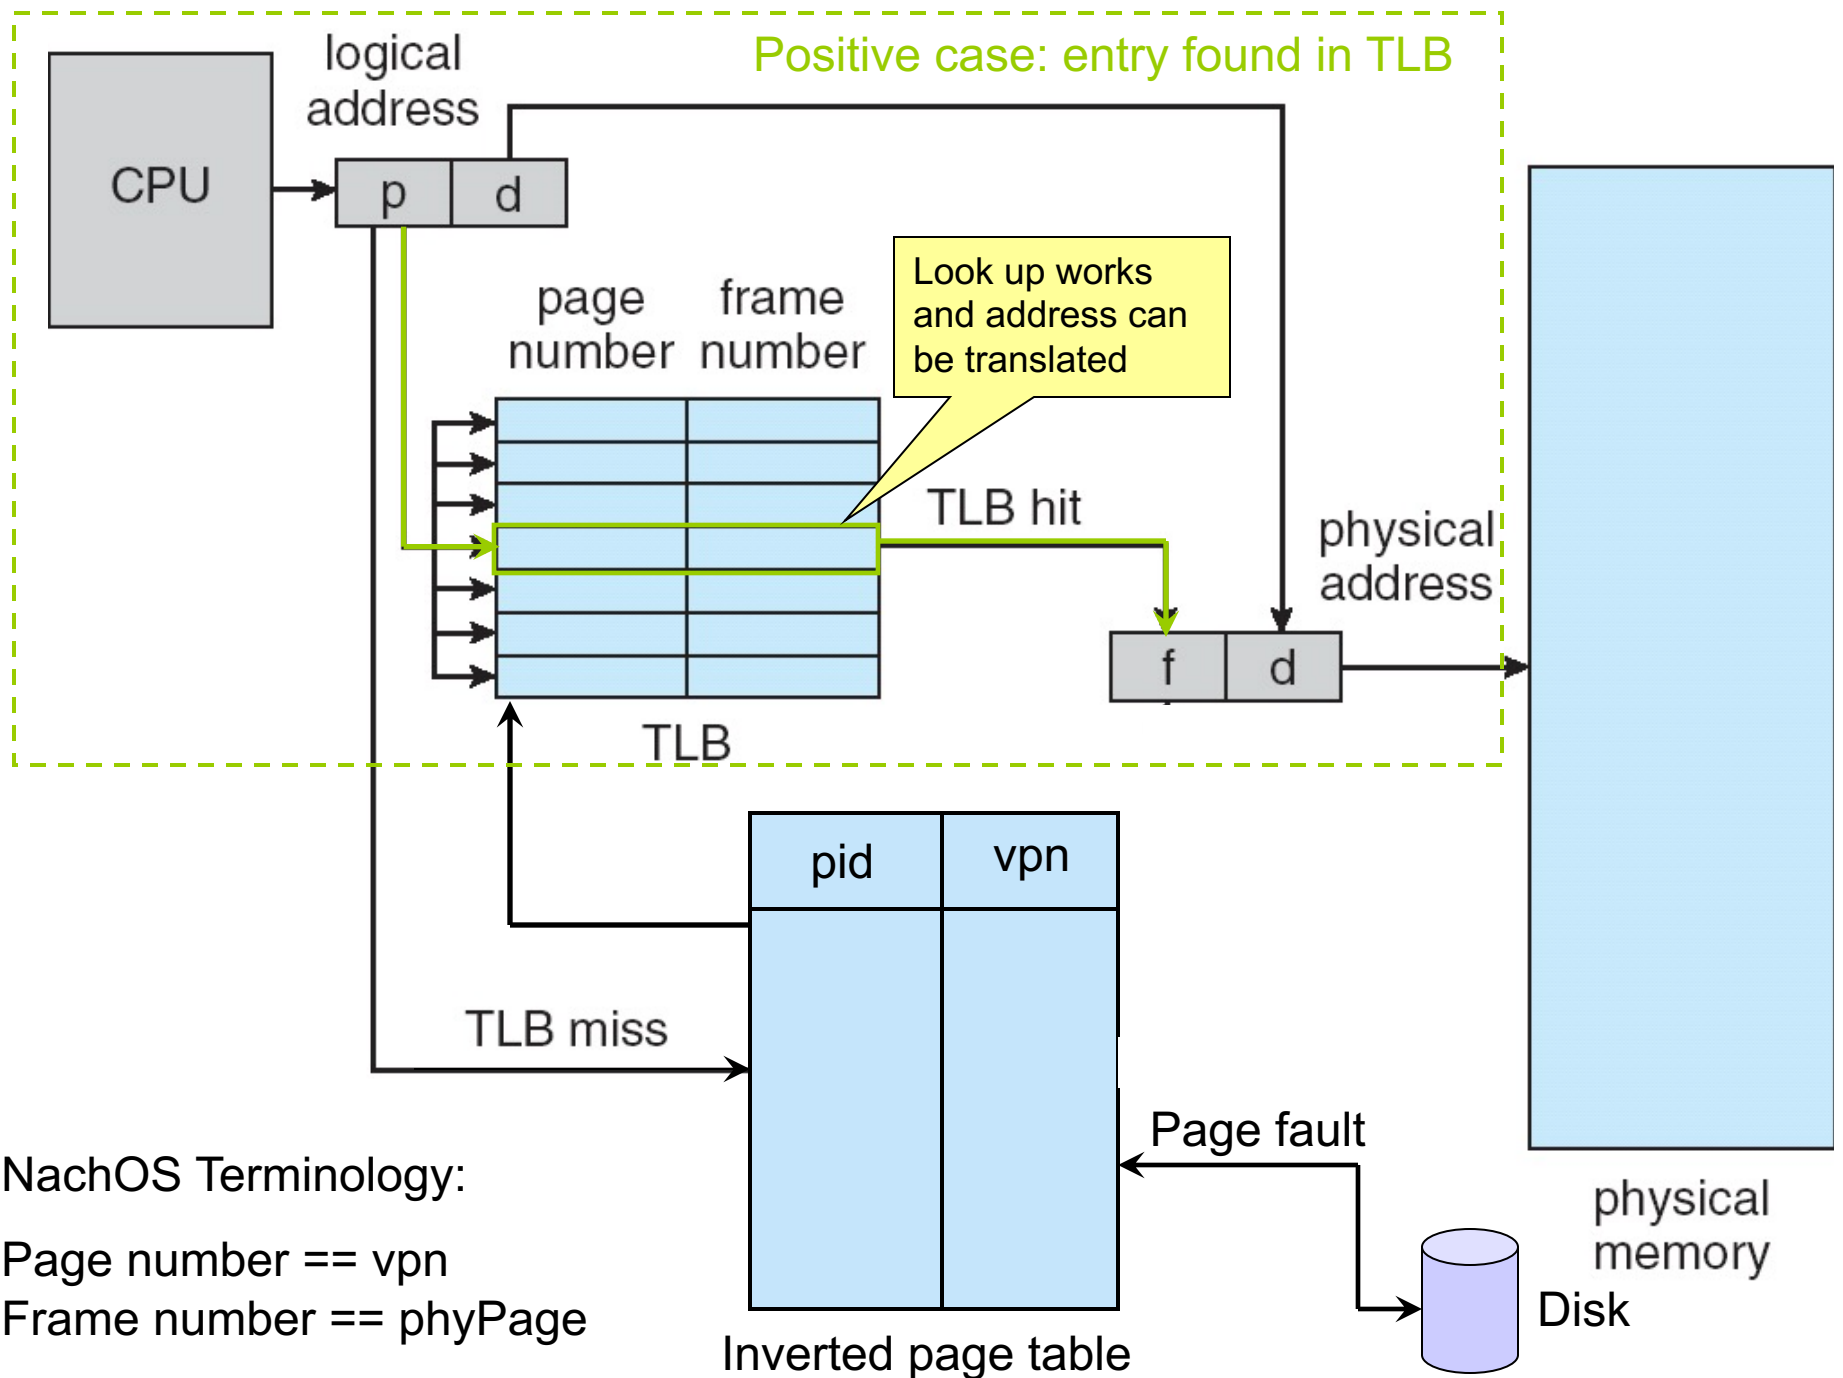
# Address Translation in NachOS

- How to translate a virtual address into a physical address?
  - Translation Look-aside Buffer (TLB): one per machine
  - Inverted Page Table (IPT): one for the entire system
- Call `Translate(vpn, &phyAddr ...)`  
(defined under `machine/translate.cc`)
  - Calculate VPN (Virtual Page Number)
  - Lookup VPN in TLB
  - If lookup was successful
    - Calculate physical address
  - If lookup was not successful (i.e., TLB miss)
    - Generate an exception
- If an exception is generated
  - Call `updateTLB(vpn)`
  - Call `Translate(vpn, &phyAddr ...)` again

# Virtual Address Translation

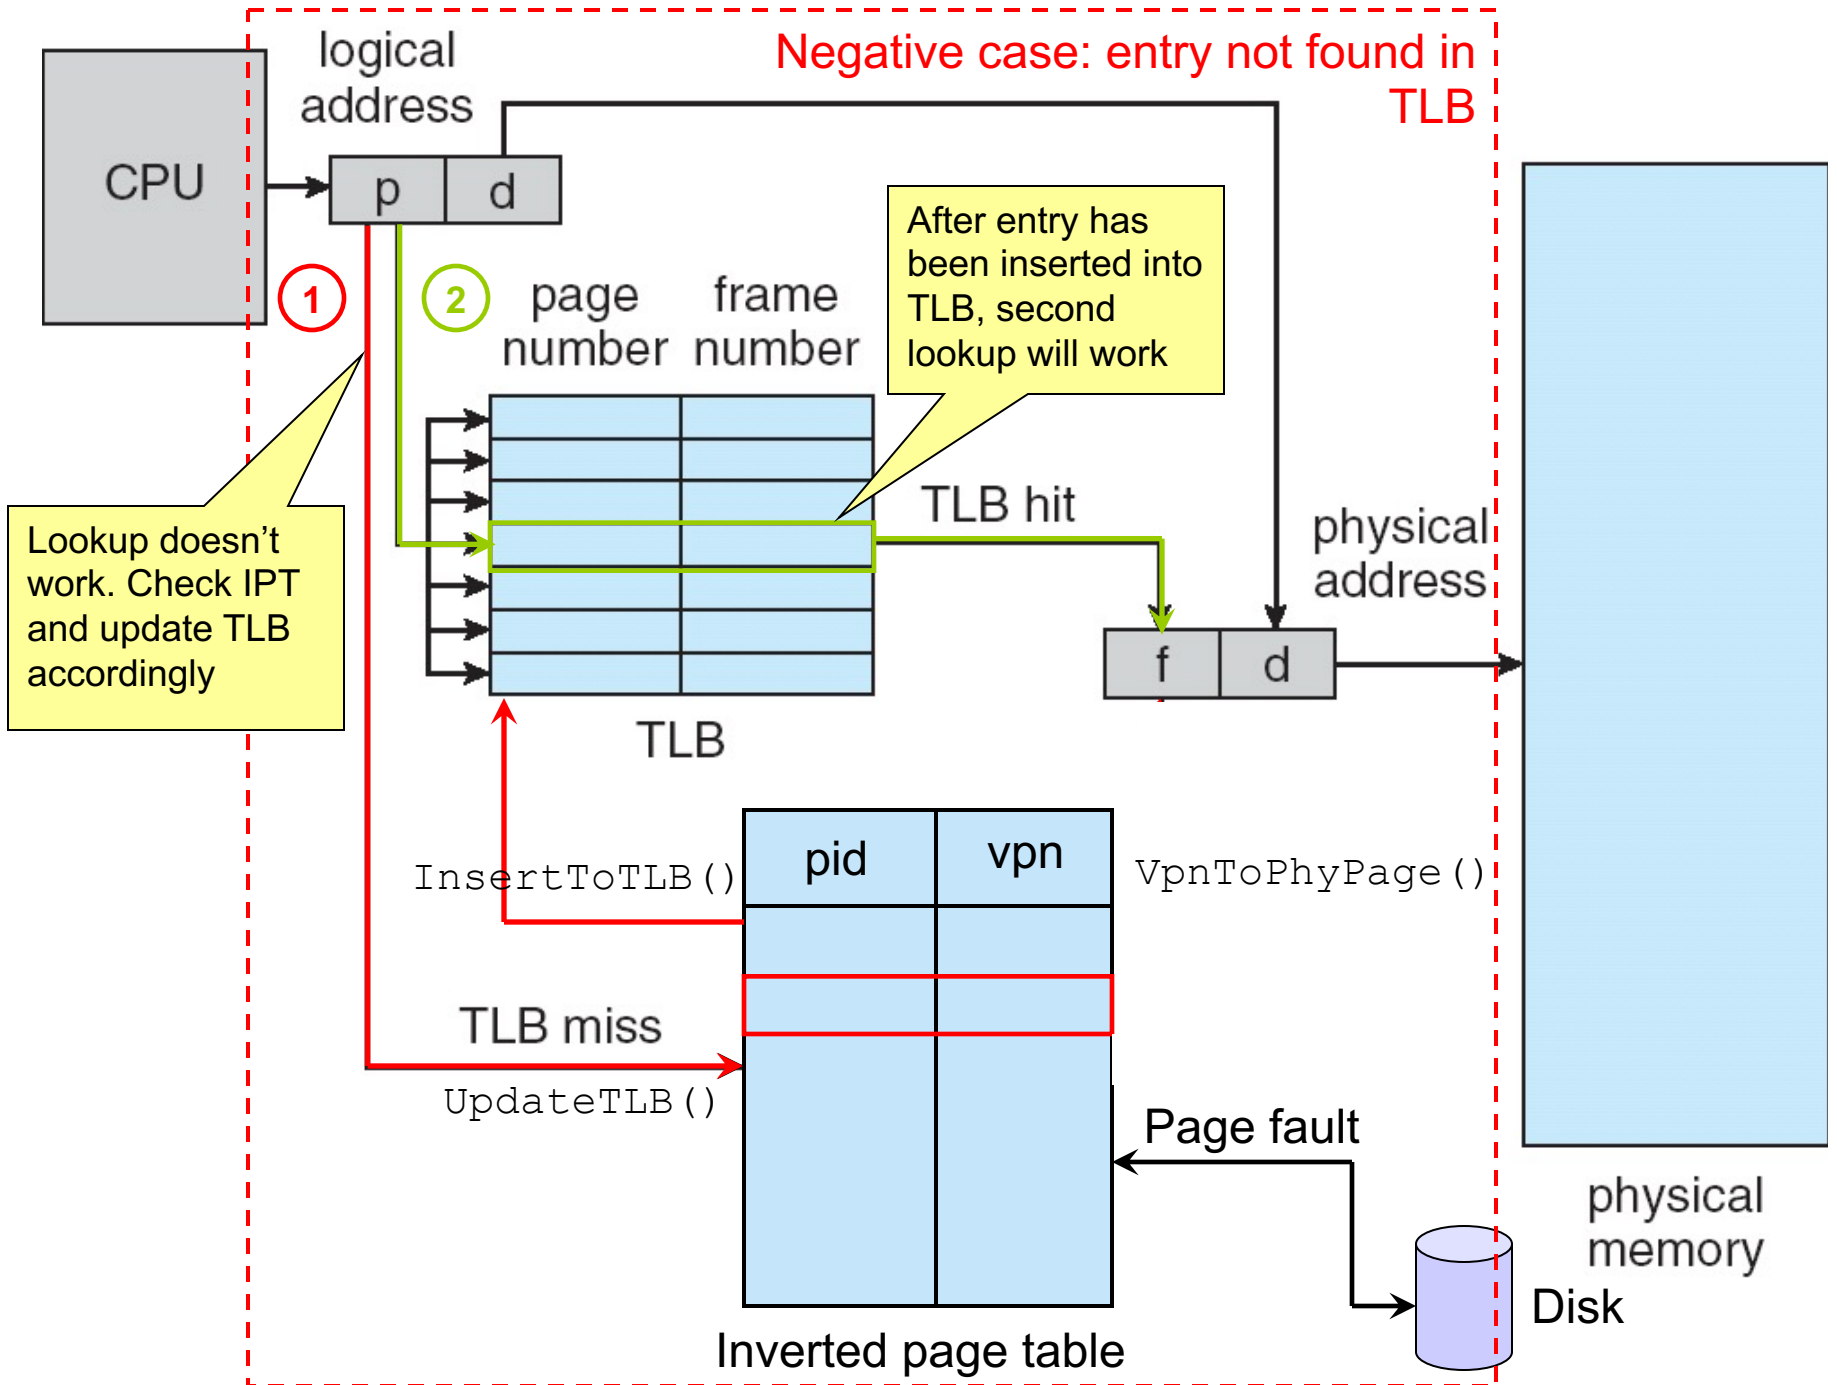






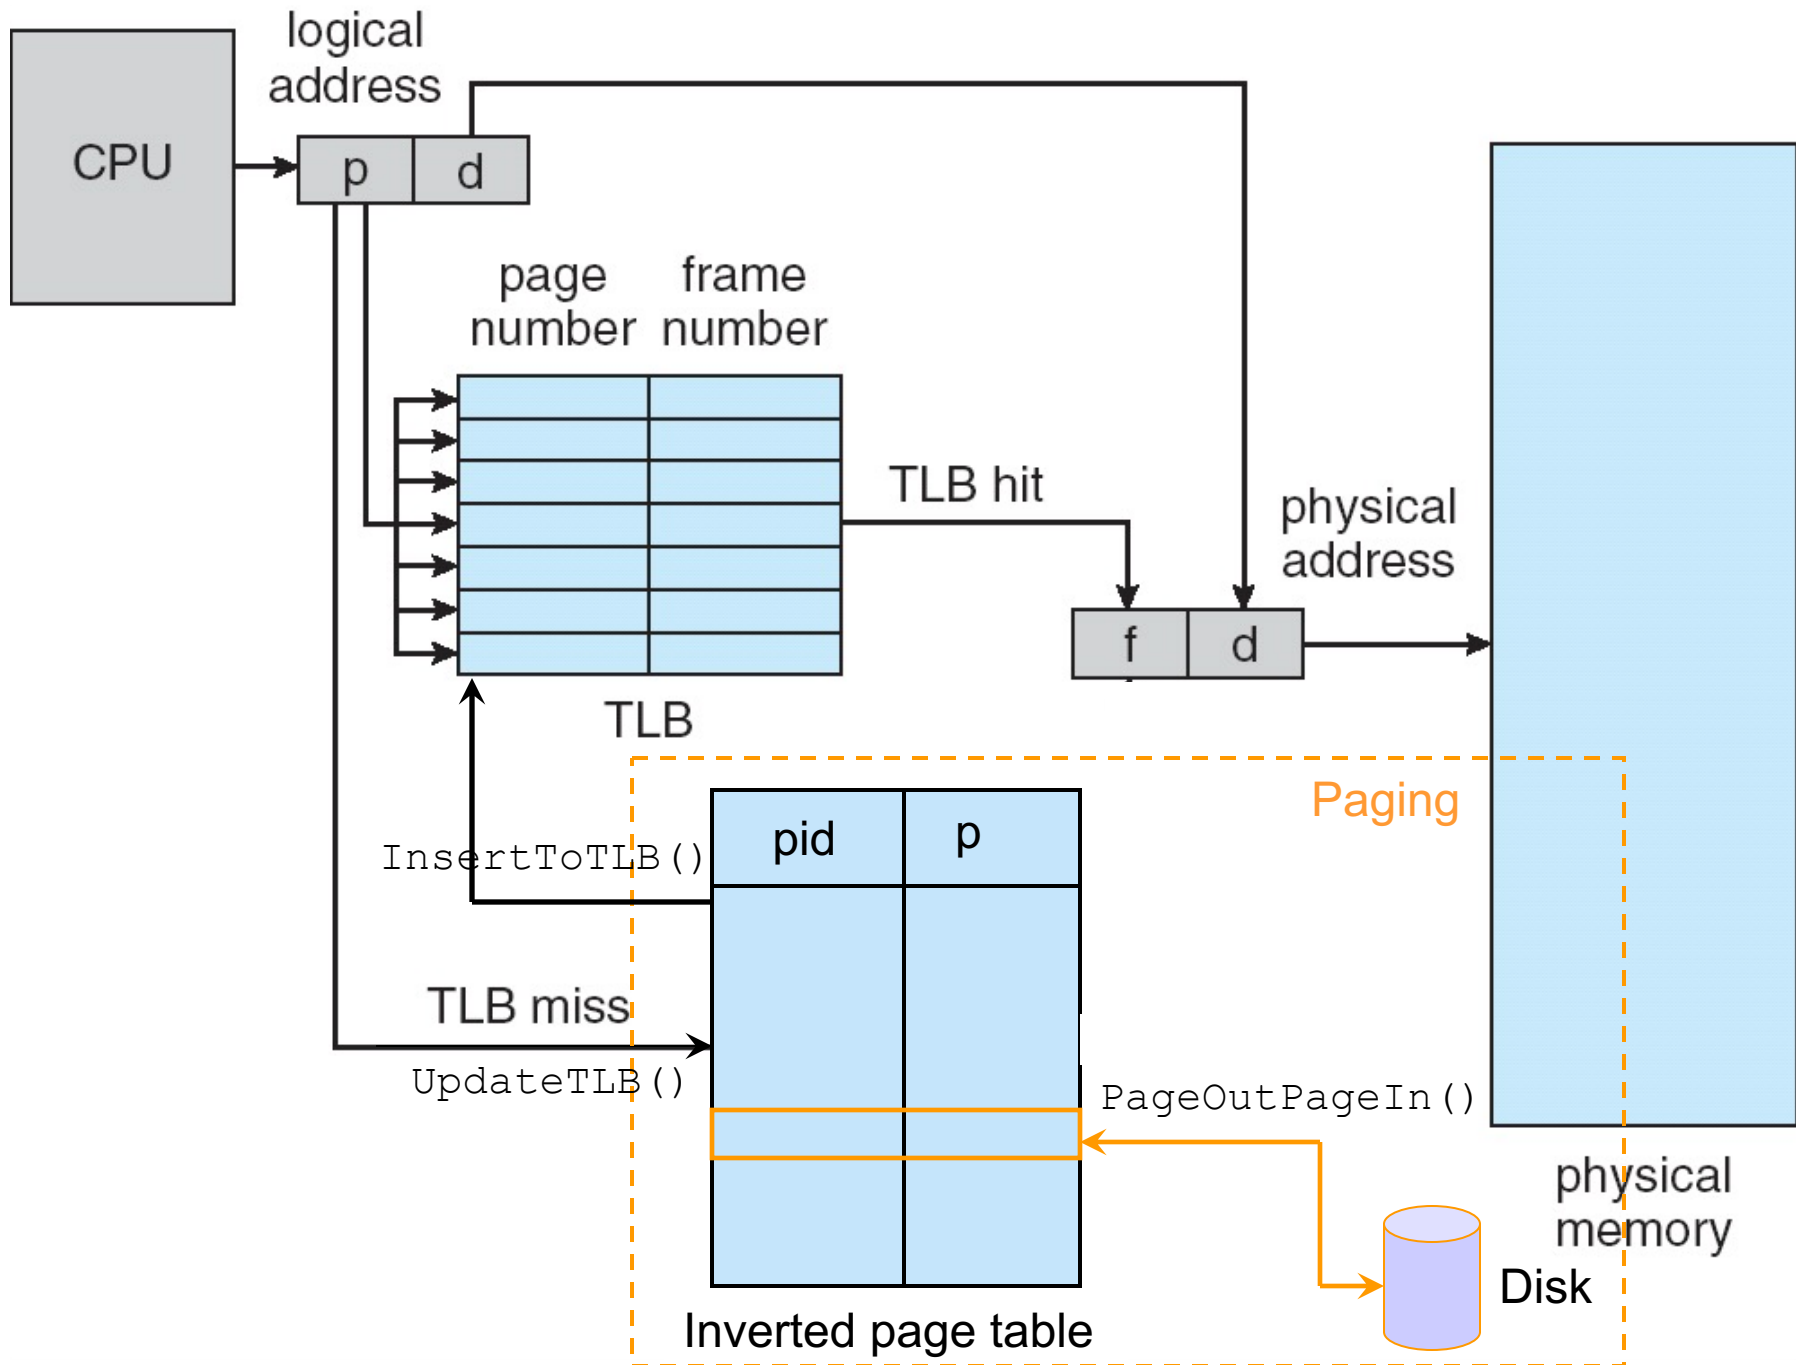
# Experiment 4 – Update TLB

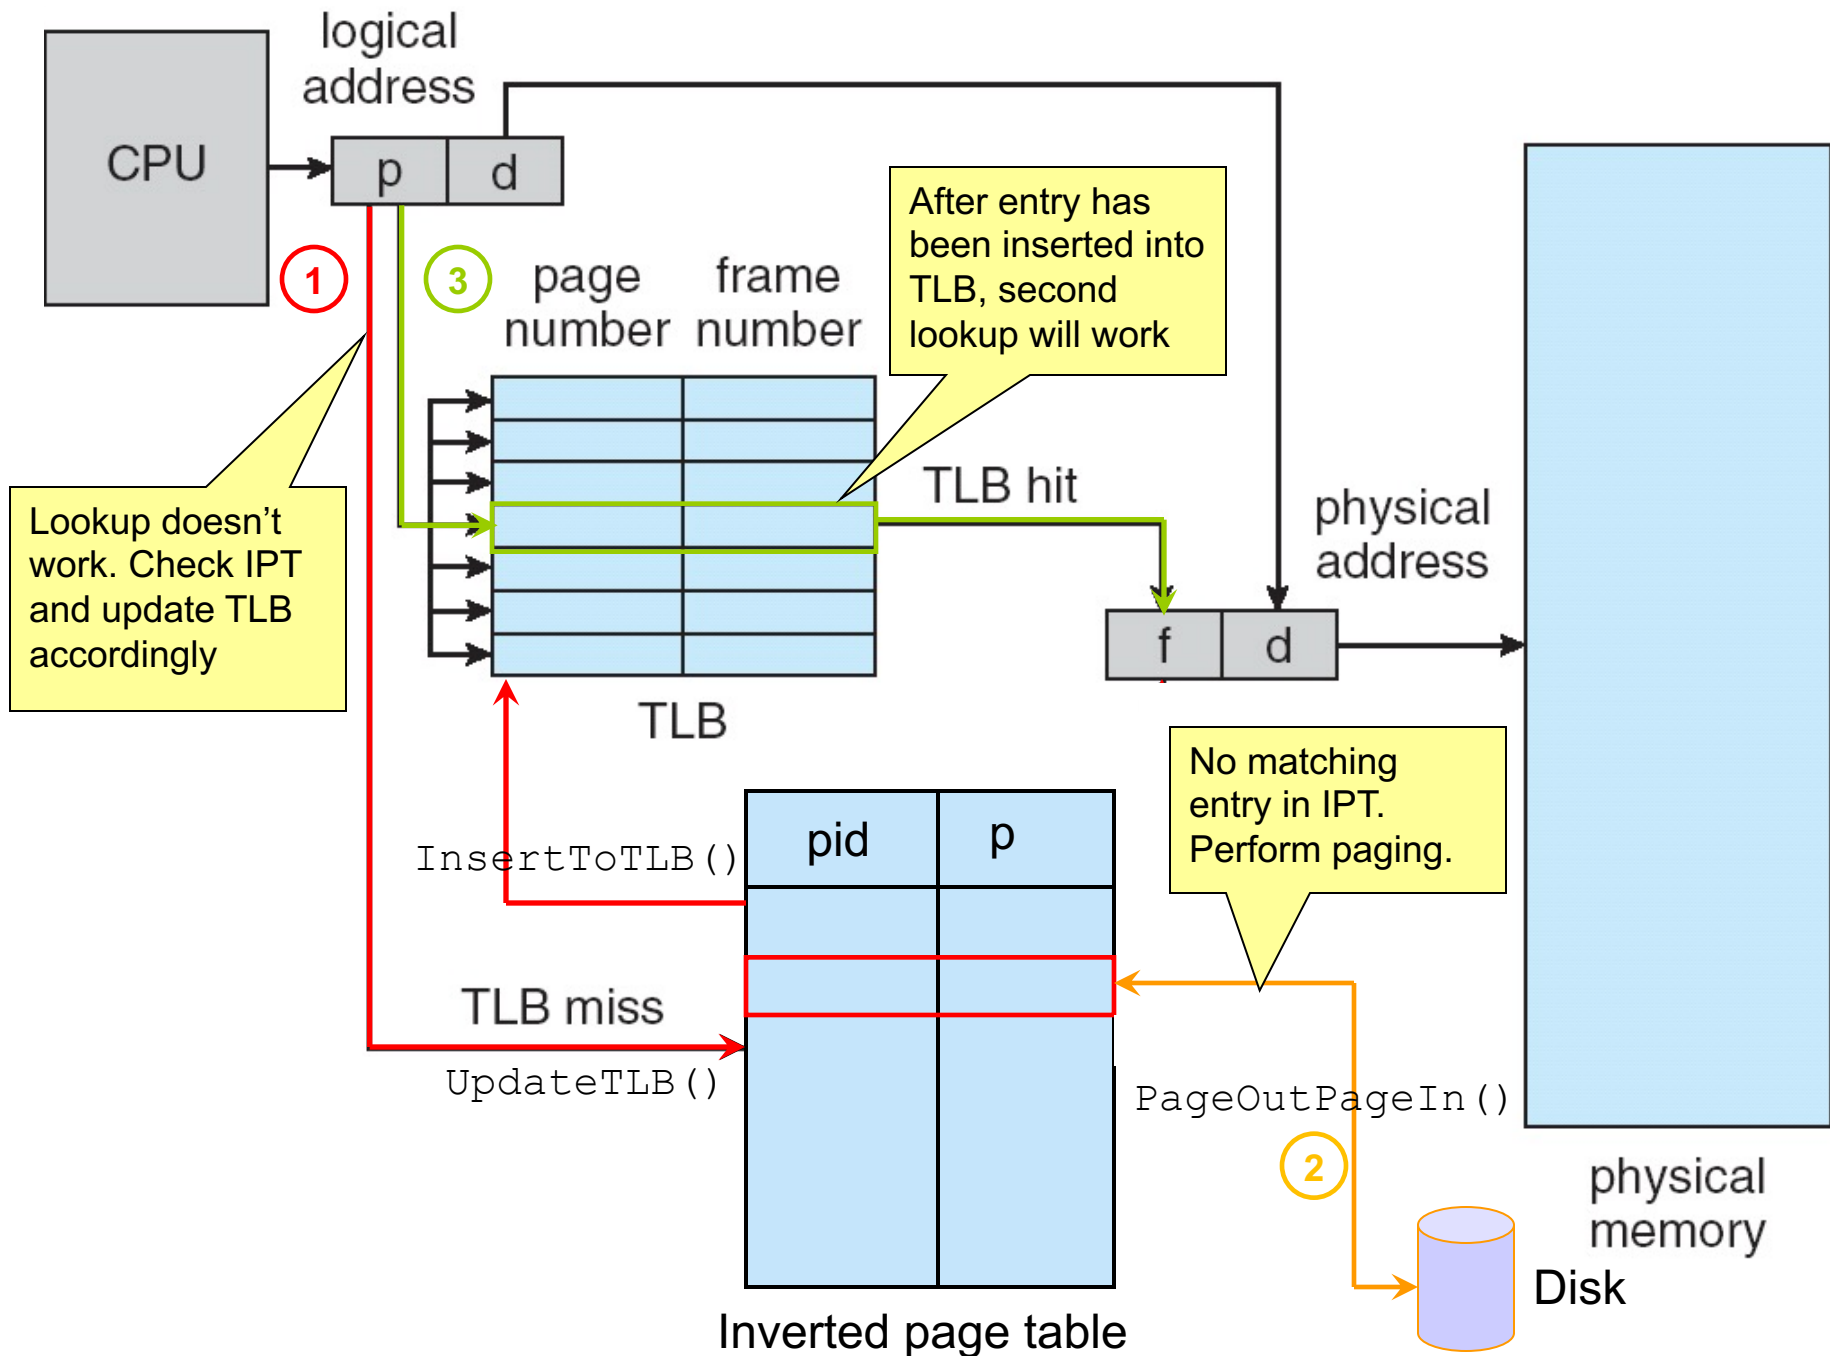
- In case the lookup was not successful (i.e., in case of a TLB miss), how to update the TLB?
- **Function: `UpdateTLB(vpn)`**  
**(defined under `vm/tlb.cc`)**
  - Call `vpnToPhyPage(vpn)`
  - If `vpn` can be found in IPT (i.e., `vpnToPhyPage(vpn)` returns a valid `phyPage`), update TLB
    - Call `insertToTLB(vpn, phyPage)`
  - Otherwise, perform paging and update TLB
    - Call `insertToTLB(vpn, PageOutPageIn(vpn))`

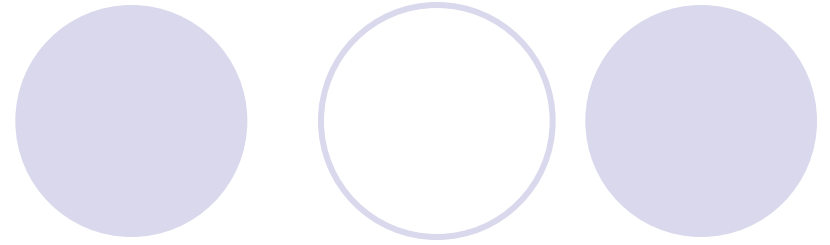
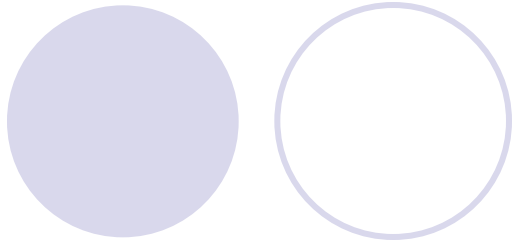


# Experiment 4 – Update TLB

- If VPN cannot be found in IPT, how to perform paging?
- **Function:** `PageOutPageIn (vpn)`
  - Determine the victim frame using the least recently used algorithm
  - Page out victim page
    - Write victim page to swap file
  - Page in the new page
    - Load new physical page from swap file
    - Update the IPT table with VPN/PhyPage combination







## Discussion of Experiment 4

# Experiment 4 – Overview

- Objective

- Understand why TLB can provide fast address translation
- Know how address translation is done by using IPT
- Understand how page replacement is essential to virtual memory
- Understand how to implement a least recently used replacement algorithm

- Tasks

- Implement missing functionality for address translation
- Run test program and analyse output



# Directory Structure

<b>bin</b>	For generating NachOS format files, <b>DO NOT CHANGE!</b>
<b>filesystem</b>	NachOS kernel related to file system, <b>DO NOT CHANGE!</b>
<b>exp3</b>	<a href="#">Experiment 3, process synchronization.</a>
<b>machine</b>	MIPS H/W simulation, <b>DO NOT CHANGE</b> unless asked.
Makefile.common	For compilation of NachOS,
Makefile.dep	<b>DO NOT CHANGE!</b>
<b>network</b>	NachOS kernel related to network, <b>DO NOT CHANGE!</b>
<b>port</b>	NachOS kernel related to port, <b>DO NOT CHANGE!</b>
readme	Short description of OS labs and assessments
<b>test</b>	NachOS format files for testing virtual memory, <b>DO NOT CHANGE!</b>
<b>threads</b>	NachOS kernel related to thread management, <b>DO NOT CHANGE!</b>
<b>userprog</b>	NachOS kernel related to running user applications, <b>DO NOT CHANGE!</b>
<b>vm</b>	<a href="#">Experiment 4, coding virtual memory (TLB, page replacement)</a>

# Experiment 4 – Overview

- Incomplete VM code can be found in
  - `vm/tlb.cc`
- You need to add your code to these 3 functions:
  - `int VpnToPhyPage(int vpn)`
  - `void InsertToTLB(int vpn, int phyPage)`
  - `int lruAlgorithm(void)`
- Binary test program for Experiment 4 is provided
  - Execute the test program:
    - `./nachos -x ../test/vmtest.noff -d`

# Experiment 4 – VpnToPhyPage()

- Return the physical page for a VPN (if it exists in the IPT)
  - `int VpnToPhyPage(int vpn) { ... }`
- IPT is realised as memory table (see `vm/ipt.h`).
- A memory table is used by the algorithm
  - Can be accessed by using `memoryTable[i]`
- This memory table has as many entries as there are physical pages
  - A constant is used for that purpose: `NumPhysPages`
- Iterate the memory table to find the corresponding physical page entry.
- Return the index `i` for the physical page entry for which the following conditions are all true:
  - `memoryTable[i].valid`
  - `memoryTable[i].pid == currentThread->pid`
  - `memoryTable[i].vPage == vpn`
- Return -1 if no entry can be found that matches the above condition

```
//-----  
// VpnToPhyPage  
//     Gets a phyPage for a vpn, if exists in ipt.  
//-----  
  
int VpnToPhyPage(int vpn)  
{  
    //your code here to get a physical frame for page vpn  
    //you can refer to PageOutPageIn(int vpn) to see how an entry was created in ipt  
    //your code here  
    for(.....){  
        if(.....){  
            .....;  
        }  
    }  
    // return sth if conditios are not fulfilled  
    .....|;  
}  
  
//-----
```

# Experiment 4 – InsertToTLB()

- Put a virtual page number and its associated physical page into the TLB

```
void InsertToTLB(int vpn, int phyPage) {  
    int i = 0; //entry in the TLB  
  
    //your code to find an empty in TLB or to  
    //replace the oldest entry if TLB is full  
    ???  
  
    //copy dirty data to memoryTable  
    ...  
    memoryTable[phyPage].lastUsed = stats->totalTicks;  
}
```

This is the set to the current tick, mentioned in the lab manual. It's already done for you, so you don't need to worry about it

- What do you need to do?
  - Replace the ??? with your code, don't touch the rest
  - Your code has to make sure that *i* is set correctly
    - i.e., the *i*-th entry in the TLB is either *empty* or the *oldest*

# Experiment 4 – InsertToTLB()

- Check all entries in the TLB whether there are any invalid entries
- How to do that?
  - The TLB is an array of `TranslationEntry` objects:
    - `TranslationEntry *tlb;`
    - You can access the TLB in the following way:
      - `machine->tlb[i]`
  - A translation entry has several flags
    - For example: you can check whether an entry is valid:
      - `machine->tlb[i].valid`
    - Check `machine/translate.h` for more details
  - The size of that array is defined by the constant `TLBSize`

# Experiment 4 – InsertToTLB()

- If there is an invalid TLB entry, then  $i$  should point to it
- For example, if the 2<sup>nd</sup> entry is invalid then  $i=1$ 
  - The new VPN/PhyPage value will be inserted into the 2<sup>nd</sup> entry of the table

# Entry	VPN	PhyPage	Valid	ReadOnly	Use	Dirty
0	...	...	TRUE	...	...	...
1	...	...	FALSE	...	...	...
2	...	...	TRUE	...	...	...

# Experiment 4 – InsertToTLB()

- If there is no invalid entry, then  $i$  should point to the oldest entry
  - The oldest entry will be replaced by the new VPN/PhyPage values
  - You need to keep track of the oldest entry
  - C++ hint: use a static variable for that purpose
    - `static int FIFOPointer = 0;`
    - Once a static variable is initialised, it remains in the memory
    - No re-initialisation afterwards
  - Make sure FIFOPointer is always correctly pointing to the oldest entry
    - Simple FIFO: if an entry is just inserted, then the entry next to it is the oldest entry
    - `FIFOPointer = (i + 1) % TLBSize`



```

//-----
// InsertToTLB
// Put a vpn/phyPage combination into the TLB. If TLB is full, use FIFO
// replacement
//-----

void InsertToTLB(int vpn, int phyPage)
{
    int i = 0; //entry in the TLB


    //your code to find an empty in TLB or to replace the oldest entry if TLB is full
    //declare FIFOPointer here

    // find an empty in TLB
    for (.....)
        if (.....) {
            .....;
            break;
        }

    //replace the oldest entry if TLB is full
    if (.....) {
        .....;
        .....;
    }
}

```

Print out TLB info to fill in Table1



```

// print out TLB
printf("TLB:\t");
for (int j = 0; j < TLBSize; j++){
    printf("%i,%i,%i%c\t", machine->tlb[j].virtualPage, machine->tlb[j].physicalPage, machine->tlb[j].valid, j == i ? '*' : ' ');
}
printf("\n");

```

```

// copy dirty data to memoryTable
if(machine->tlb[i].valid){
    memoryTable[machine->tlb[i].physicalPage].dirty=machine->tlb[i].dirty;
    memoryTable[machine->tlb[i].physicalPage].TLBentry=-1;
}

```

```

//update the TLB entry
machine->tlb[i].virtualPage = vpn;
machine->tlb[i].physicalPage = phyPage;
machine->tlb[i].valid
    = TRUE;
machine->tlb[i].readOnly
    = FALSE;
machine->tlb[i].use
    = FALSE;
machine->tlb[i].dirty
    = memoryTable[phyPage].dirty;

```

# Experiment 4 – lruAlgorithm()

- Determines which physical page should be paged out
  - `int lruAlgorithm(void)`
- Need to find the least recently used entry in the `memoryTable`.
- The last tick that the physical page is accessed, is stored in `memoryTable[i].lastUsed`.
- Search for the invalid entry from the beginning of the `memoryTable`. If there is an invalid entry, return that to be used by the virtual page.
  - `!memoryTable[i].valid`
- Otherwise, find a victim (entry with smallest `lastUsed`)
  - Return the page number

# lruAlgorithm

```
|  
//-----  
// lruAlgorithm  
// Determine where a vpn should go in phymem, and therefore what  
// should be paged out. This lru algorithm is the one discussed in the  
// lectures.  
//-----  
  
int lruAlgorithm(void)  
{  
    //your code here to find the physical frame that should be freed  
    //according to the LRU algorithm.  
    int phyPage = 0;  
  
    for(.....){  
        if(.....){  
            ....;  
        }  
  
        if(.....){  
            .....;  
        }  
    }  
  
    return phyPage;  
}  
//
```

# PageOutPageIn

```
//-----  
// PageOutPageIn  
//      Calls DoPageOut and DoPageIn and handles memoryTable  
// bookkeeping. Use lru algorithm to find the replacement page.  
//-----
```

```
int PageOutPageIn(int vpn)  
{  
    int phyPage;  
  
    //increase the number of page faults  
    stats->numPageFaults++;  
    //call the LRU algorithm, which returns the freed physical frame  
    phyPage=lruAlgorithm();
```

Print out IPT info to fill into Table1

```
|// print out memory table  
|for (int i = 0; i < NumPhysPages; i++){  
|    printf("%i,%i,%i,%i%c\t", memoryTable[i].pid, memoryTable[i].vPage, memoryTable[i].lastUsed, memoryTable[i].valid, phyPage == i ? '*' : ' ');  
|}  
|printf("\n");
```

```
//Page out the victim page to free the physical frame  
DoPageOut(phyPage);  
//Page in the new page to the freed physical frame  
DoPageIn(vpn, phyPage);
```

```
//update memoryTable for this frame  
memoryTable[phyPage].valid=TRUE;  
memoryTable[phyPage].pid=currentThread->pid;  
memoryTable[phyPage].vPage=vpn;  
memoryTable[phyPage].dirty=FALSE;  
memoryTable[phyPage].TLBentry=-1;  
memoryTable[phyPage].lastUsed=0;  
memoryTable[phyPage].swapPtr=currentThread->space->swapPtr;
```

```
return phyPage;
```

```
}
```

# Experiment 4 – Analysis of Output

tick	vpn	pid	IPT[0]	IPT[1]	IPT[2]	IPT[3]	TLB[0]	TLB[1]	TLB[2]	Page Out
			pid, vpn, last used, valid				vpn, phy, valid			
10	0	0	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0	0,0,0	0,0,0	N
13	9	0	0,0,12,1	0,0,0,0	0,0,0,0	0,0,0,0	0,0,1	0,0,0	0,0,0	N
15	26	0	0,0,12,1	0,9,15,1	0,0,0,0	0,0,0,0	0,0,1	9,1,1	0,0,0	N
20	1	0	0,0,12,1	0,9,19,1	0,26,17,1	0,0,0,0	0,0,1	9,1,1	26,2,1	N

- Analyse the output of your program and complete **Table1.csv** (record down the entries of the IPT and TLB **before** replacement and **highlight** the entry that is selected to be updated in **Table1.pdf**. **Do not highlight cells in Table1.csv and do not add any new column header to Table1.csv**)
  - For IPT entries, record:
    - pid, vpn, last used, valid
  - For TLB entries, record:
    - vpn, phy, valid
- Write down the following (By completing **Table2.csv**)
  - Page size (defined in NachOS)
  - Number of physical frames (defined in NachOS)
  - TLB size (defined in NachOS)
  - Number of pages used by the test program
  - Number of page faults that occurred during execution of the test program

# Table1.csv

- Do not highlight cells in Table1.csv and do not add any new column header to Table1.csv
- VERY IMPORTANT**
- Replace Table1.pdf submission

tick	vpn	pid	IPT[0]	IPT[1]	IPT[2]	IPT[3]	TLB[0]	TLB[1]	TLB[2]	Page Out
			pid, vpn, last used, valid				vpn, phy, valid			
10	0	0	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0	0,0,0	0,0,0	N
13	9	0	0,0,12,1	0,0,0,0	0,0,0,0	0,0,0,0	0,0,1	0,0,0	0,0,0	N
15	26	0	0,0,12,1	0,9,15,1	0,0,0,0	0,0,0,0	0,0,1	9,1,1	0,0,0	N
20	1	0	0,0,12,1	0,9,19,1	0,26,17,1	0,0,0,0	0,0,1	9,1,1	26,2,1	N

- By add \* to mark the entry that is selected to be updated.**
- Example:**

10	0	0	0,0,0,0*	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0*	0,0,0	0,0,0	N
13	9	0	0,0,12,1	0,0,0,0*	0,0,0,0	0,0,0,0	0,0,1	0,0,0*	0,0,0	N
15	26	0	0,0,12,1	0,9,15,1	0,0,0,0*	0,0,0,0	0,0,1	9,1,1	0,0,0*	N
20	1	0	0,0,12,1	0,9,19,1	0,26,17,1	0,0,0,0*	0,0,1*	9,1,1	26,2,1	N

# Tabulated table1 based on output.txt

10	0	0 0,0,0*	0,0,0,0	0,0,0,0	0,0,0,0	0,0,0*	0,0,0	0,0,0	N
13	9	0 0,0,12,1	0,0,0,0*	0,0,0,0	0,0,0,0	0,0,1	0,0,0*	0,0,0	N
15	26	0 0,0,12,1	0,9,15,1	0,0,0,0*	0,0,0,0	0,0,1	9,1,1	0,0,0*	N
20	1	0 0,0,12,1	0,9,19,1	0,26,17,1	0,0,0,0*	0,0,1*	9,1,1	26,2,1	N

== Tick 10 ==

interrupts: on -> off

Time: 10, interrupts off

Pending interrupts:

In mapcar, about to invoke 804b454(8752098)

Interrupt handler timer, scheduled at 96

End of pending interrupts

interrupts: off -> on

Scheduling interrupt handler the console read at time = 110

Initializing address space, num pages 27, size 3456

Initializing stack register to 3440

Starting thread "main" at time 10

Reading VA 0x0, size 4

Translate 0x0, read: \*\*\* no valid TLB entry found for this virtual page 0!

Exception: page fault/no TLB entry

0,0,0,0\*0,0,0,0 0,0,0,0 0,0,0,0

paging in: pid 0. phyPage 0. vpn 0

TLB: 0,0,0\* 0,0,0 0,0,0

The corresponding TLBentry for Page 0 in TLB is 0

# Table2.csv

Can be found in certain header  
file belonging to machine directory  
(see exp4 detailed handout)

PageSize
Number of physical frames
TLB size
Ticks
Paging Faults
Paging Out
TLB Miss

From output.txt,

Ticks: .... **Copied all tick info into table**  
Disk I/O: reads 0, writes 0  
Console I/O: reads 0, writes 0  
Paging: .....|  
Network I/O: packets received 0, sent 0  
  
Cleaning up...



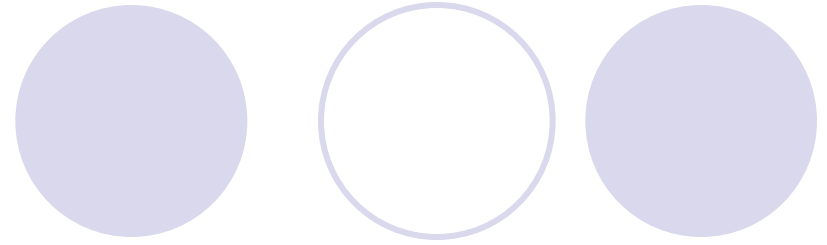
# Experiment 4 – Summary

- No group effort
- **Assessment:**
  - Assessment of your implementation. Please leave your code, the output file **output.txt** as well as **Table1.csv**, **Table2.csv** and ~~**Table1.pdf**~~ files containing the above information in the **vm** folder for TA/Supervisor to review. **Deadline is 1 week after your lab session (e.g., if lab session is from 10AM-12PM on a Monday, then deadline is 9:59AM on the next Monday).**
  - **Lab Quiz 2**, which is an online multiple-choice quiz, will be administered through NTULearn in Week 14.

## Updated instruction

- In the vm folder, you should have **the code (including the compiled nachos), output.txt, Table1.csv (using \* to mark entry to be updated), Table2.csv.**
- **Table1.pdf is not needed.**

# Acknowledgement



- The slides are revised from the previous versions created by Dr. Heiko Aydt and Tan Wen Jun.